

MIT-BIH noise stress test database

October 18, 2025

```
[31]: # Step 1: Import libraries
import os
import wfdb
import pandoc
import zipfile
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.inspection import permutation_importance
from scipy.signal import find_peaks
from scipy.stats import skew, kurtosis
from sklearn.impute import SimpleImputer
from imblearn.over_sampling import SMOTE
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.neural_network import MLPClassifier
from sklearn.metrics import accuracy_score, precision_score, recall_score, \
    f1_score, roc_auc_score, confusion_matrix, roc_curve, auc, \
    classification_report

[3]: base_path = "/Users/mohithreddy/mit-bih-noise-stress-test-database-1.0.0/"
    mit_bih_noise_stress_test_database_1_0_0 = base_path + "mit-bih-noise-stress-test-database-1.0.0"

[4]: zip_path = "/Users/mohithreddy/Downloads/mit-bih-noise-stress-test-database-1.0.0.0.zip" # change if needed
    extract_dir = "/Users/mohithreddy/mit-bih-noise-stress-test-database-1.0.0"

    # Extract the zip
    with zipfile.ZipFile(zip_path, 'r') as zip_ref:
        zip_ref.extractall(extract_dir)

    print("Extracted to:", extract_dir)
    print("Files inside:", os.listdir(extract_dir)[:10])
```

Extracted to: /Users/mohithreddy/mit-bih-noise-stress-test-database-1.0.0
Files inside: ['mit-bih-noise-stress-test-database-1.0.0', '__MACOSX']

```

[5]: os.getcwd()

[5]: '/Users/mohithreddy'

[6]: os.listdir("mit-bih-noise-stress-test-database-1.0.0")[:10]

[6]: ['mit-bih-noise-stress-test-database-1.0.0', '__MACOSX']

[7]: data_path = "mit-bih-noise-stress-test-database-1.0.0/
↳ mit-bih-noise-stress-test-database-1.0.0"

[8]: os.listdir(data_path)[:10]

[8]: ['119e12.he',
      '119e06.he',
      'bw.he',
      'em.he-',
      'em.xws',
      'ma.dat',
      '118e12.xws',
      '118e06.xws',
      'RECORDS',
      '118e24.atr']

[9]: data_path = "/Users/mohithreddy/Downloads/mit-bih-noise-stress-test-database-1.
↳ 0.0/mit-bih-noise-stress-test-database-1.0.0"

[10]: # Show available header files
hea_files = [f for f in os.listdir(extract_dir) if f.endswith('.hea')]
print("Total .hea files:", len(hea_files))
print(hea_files[:15]) # show first 15

Total .hea files: 0
[]

[11]: data_path = "mit-bih-noise-stress-test-database-1.0.0/
↳ mit-bih-noise-stress-test-database-1.0.0"

# Select 5 subjects
subjects = ['118e_6', '118e12', '118e18', '119e_6', '119e24']

records = {}

for subj in subjects:
    record_path = os.path.join(data_path, subj)
    record = wfdb.rdrecord(record_path)
    annotation = wfdb.rdann(record_path, 'atr') # apnea annotations

```

```

records[subj] = {
    "signal": record.p_signal[:, 0],
    "fs": record.fs,
    "annotations": annotation.sample,
    "labels": annotation.symbol
}

print(f"{subj}: Loaded with {len(record.p_signal)} samples at {record.fs} Hz.")

```

```

118e_6: Loaded with 650000 samples at 360 Hz.
118e12: Loaded with 650000 samples at 360 Hz.
118e18: Loaded with 650000 samples at 360 Hz.
119e_6: Loaded with 650000 samples at 360 Hz.
119e24: Loaded with 650000 samples at 360 Hz.

```

```

[12]: def label_class(ann):
    """
    Map annotation symbols to binary labels:
    Normal (0) and Abnormal (1)
    """
    # Normal beat types according to AAMI standards
    normal_symbols = ['N', 'L', 'R', 'e', 'j']

    labels = []
    for symbol in ann.symbol:
        if symbol in normal_symbols:
            labels.append(0) # Normal
        else:
            labels.append(1) # Abnormal
    return np.array(labels)

```

```

[13]: def label_class_from_symbols(symbols):
    """Convert annotation symbols to binary labels: 0=Normal, 1=Abnormal"""
    normal_symbols = ['N', 'L', 'R', 'e', 'j']
    return np.array([0 if sym in normal_symbols else 1 for sym in symbols])

def create_window_targets(records, window_size_sec=3, overlap_sec=0.5):
    label_data = []

    for subj, rec in records.items():
        signal = rec["signal"]
        fs = rec["fs"]

        # --- handle both dictionary styles safely ---
        if "annotation" in rec:
            ann = rec["annotation"]
            beat_symbols = ann.symbol

```

```

        beat_positions = np.array(ann.sample)
    elif "annotations" in rec and "labels" in rec:
        # old format from your earlier code
        beat_symbols = rec["labels"]
        beat_positions = np.array(rec["annotations"])
    else:
        raise KeyError(f" Subject {subj} has no annotation data!")

    # Convert beat types to 0/1 labels
    beat_labels = label_class_from_symbols(beat_symbols)

    # --- create windows ---
    win_len = int(fs * window_size_sec)
    step = int(fs * (window_size_sec - overlap_sec))

    for start in range(0, len(signal) - win_len, step):
        end = start + win_len
        beats_in_window = beat_labels[(beat_positions >= start) &
        ↪(beat_positions < end)]
        target = 1 if np.any(beats_in_window == 1) else 0
        label_data.append({
            "Subject_ID": subj,
            "Start": start,
            "End": end,
            "Target": target
        })

    print(f" Processed {subj}: {len(label_data)} windows labeled so far")

    return pd.DataFrame(label_data)

# --- Run the function ---
target_df = create_window_targets(records, window_size_sec=3, overlap_sec=0.5)

print("\n Target labels created successfully!")
print(target_df["Target"].value_counts())

```

```

Processed 118e_6: 722 windows labeled so far
Processed 118e12: 1444 windows labeled so far
Processed 118e18: 2166 windows labeled so far
Processed 119e_6: 2888 windows labeled so far
Processed 119e24: 3610 windows labeled so far

```

```

Target labels created successfully!
Target
0      2232
1      1378

```

Name: count, dtype: int64

```
[14]: # --- Define the labeling function if not already done ---
normal_class = ['N']
abnormal_class = ['A', 'V', 'L', 'R', 'E']

def label_beats(annotation):
    """Return binary labels for MIT-BIH arrhythmia beats."""
    labels = []
    for sym in annotation.symbol:
        if sym in normal_class:
            labels.append(0)    # Normal
        elif sym in abnormal_class:
            labels.append(1)    # Abnormal
        else:
            labels.append(np.nan) # Ignore other types
    return np.array(labels)

# --- Check labels for all selected subjects ---
subjects = ['118e_6', '118e12', '118e18', '119e_6', '119e24']

for rec in subjects:
    ann = wfdb.rdann(os.path.join(data_path, rec), "atr") # load annotation
    # file
    labels = label_class(ann)
    unique, counts = np.unique(labels[~np.isnan(labels)], return_counts=True)
    # ignore NaN
    print(f"Subject {rec} - Label Distribution:")
    for u, c in zip(unique, counts):
        label_name = "Normal" if u == 0 else "Abnormal"
        print(f"    {label_name}: {c}")
    print("-" * 40)
```

Subject 118e_6 - Label Distribution:

Normal: 2166

Abnormal: 135

Subject 118e12 - Label Distribution:

Normal: 2166

Abnormal: 135

Subject 118e18 - Label Distribution:

Normal: 2166

Abnormal: 135

Subject 119e_6 - Label Distribution:

Normal: 1543

Abnormal: 551

Subject 119e24 - Label Distribution:

Normal: 1543

Abnormal: 551

```
[15]: def segment_signal(signal, fs, window_size_sec, overlap_sec):
```

```
    """
```

```
    Splits the ECG signal into overlapping windows.
```

```
    Parameters:
```

```
        signal (array): ECG signal array
```

```
        fs (int): Sampling frequency (Hz)
```

```
        window_size_sec (float): window length in seconds
```

```
        overlap_sec (float): overlap length in seconds
```

```
    Returns:
```

```
        windows (np.ndarray): Array of segmented windows
```

```
    """
```

```
    win_len = int(fs * window_size_sec)
```

```
    step = int(fs * (window_size_sec - overlap_sec))
```

```
    windows = []
```

```
    for start in range(0, len(signal) - win_len, step):
```

```
        windows.append(signal[start:start + win_len])
```

```
    return np.array(windows)
```

```
[16]: fs = 360 # sampling frequency
```

```
signal = records['118e_6']['signal'] # example subject
```

```
window_sizes = [1, 2, 3, 4, 5] # in seconds
```

```
overlap = 0.5 # in seconds
```

```
[17]: def segment_signal(signal, fs, window_size_sec, overlap_sec):
```

```
    win_len = int(fs * window_size_sec)
```

```
    step = int(fs * (window_size_sec - overlap_sec))
```

```
    return np.array([signal[i:i+win_len] for i in range(0, len(signal)-win_len, step)])
```

```
fs = 360
```

```
window_size = 3
```

```
overlap = 0.5
```

```
def plot_windows(signal, fs, ws, overlap, rec_id):
```

```
    plt.figure(figsize=(12, 3))
```

```
    plt.plot(signal[:fs*10], color='gray')
```

```
    step = int(fs * (ws - overlap))
```




```
    for i in range(0, int(fs*10 - fs*ws), step):
```



```

plt.axvspan(i, i+fs*ws, color='lightblue', alpha=0.3)
plt.title(f'{rec_id} - Window Segmentation ({ws}s, overlap={overlap}s)')
plt.xlabel('Samples'); plt.ylabel('Amplitude (mV)')
plt.grid(alpha=0.5); plt.show()

for rec in subjects:
    plot_windows(records[rec]['signal'], records[rec]['fs'], window_size,
    overlap, rec)

```





```
[18]: # Dictionary to store results: {subject: {window_size: count}}
window_counts_all = {}

for rec in subjects:
    signal = records[rec]['signal']
    counts = {}
    for ws in window_sizes:
        segs = segment_signal(signal, fs, ws, overlap)
        counts[ws] = len(segs)
        print(f"Subject {rec} - Window {ws}s -> {len(segs)} segments")
    window_counts_all[rec] = counts
    print("-" * 50)

# --- Convert to DataFrame for plotting ---
window_df = pd.DataFrame(window_counts_all).T # rows = subjects, cols = window_
↪ sizes
window_df
```

```
Subject 118e_6 - Window 1s -> 3610 segments
Subject 118e_6 - Window 2s -> 1203 segments
Subject 118e_6 - Window 3s -> 722 segments
Subject 118e_6 - Window 4s -> 515 segments
```


Subject 118e_6 - Window 5s -> 401 segments

Subject 118e12 - Window 1s -> 3610 segments
Subject 118e12 - Window 2s -> 1203 segments
Subject 118e12 - Window 3s -> 722 segments
Subject 118e12 - Window 4s -> 515 segments
Subject 118e12 - Window 5s -> 401 segments

Subject 118e18 - Window 1s -> 3610 segments
Subject 118e18 - Window 2s -> 1203 segments
Subject 118e18 - Window 3s -> 722 segments
Subject 118e18 - Window 4s -> 515 segments
Subject 118e18 - Window 5s -> 401 segments

Subject 119e_6 - Window 1s -> 3610 segments
Subject 119e_6 - Window 2s -> 1203 segments
Subject 119e_6 - Window 3s -> 722 segments
Subject 119e_6 - Window 4s -> 515 segments
Subject 119e_6 - Window 5s -> 401 segments

Subject 119e24 - Window 1s -> 3610 segments
Subject 119e24 - Window 2s -> 1203 segments
Subject 119e24 - Window 3s -> 722 segments
Subject 119e24 - Window 4s -> 515 segments
Subject 119e24 - Window 5s -> 401 segments

```
[18]:
```

	1	2	3	4	5
118e_6	3610	1203	722	515	401
118e12	3610	1203	722	515	401
118e18	3610	1203	722	515	401
119e_6	3610	1203	722	515	401
119e24	3610	1203	722	515	401

```
[19]: results = {}  
  
for rec in ['118e_6', '118e12', '118e18', '119e_6', '119e24']:  
    signal = records[rec]['signal']  
    segs = segment_signal(signal, fs, window_size_sec=3, overlap_sec=0.5)  
    results[rec] = segs  
    print(f"Record {rec}: {len(segs)} windows created (3s window, 0.5s_  
↳overlap)")
```

```
Record 118e_6: 722 windows created (3s window, 0.5s overlap)  
Record 118e12: 722 windows created (3s window, 0.5s overlap)  
Record 118e18: 722 windows created (3s window, 0.5s overlap)  
Record 119e_6: 722 windows created (3s window, 0.5s overlap)  
Record 119e24: 722 windows created (3s window, 0.5s overlap)
```

```
[20]: def stability_analysis(signal, fs, window_lengths, overlaps):
    stability = {}
    for ws in window_lengths:
        for ov in overlaps:
            if ov >= ws:
                # Skip if overlap >= window
                continue
            segments = segment_signal(signal, fs, ws, ov)
            means = [np.mean(s) for s in segments]
            stability[(ws, ov)] = np.std(means)
    return stability


window_lengths = [1, 2, 3, 4, 5]
overlaps = [1, 2, 5, 10]
stability = stability_analysis(signal, fs, window_lengths, overlaps)
# Convert to DataFrame
stab_df = pd.DataFrame(list(stability.items()), columns=['(window, overlap)',
↳ 'std_mean'])
stab_df.sort_values(by='std_mean').head()
```

```
[20]: (window, overlap)  std_mean
5          (5, 1)  0.214451
6          (5, 2)  0.214506
4          (4, 2)  0.217772
3          (4, 1)  0.217967
1          (3, 1)  0.221519
```

```
[21]: plt.figure(figsize=(12, 8))

# loop through each subject and plot their first window
for i, rec in enumerate(results.keys(), start=1):
    plt.subplot(len(results), 1, i) # create one subplot per subject
    plt.plot(results[rec][0], color='black', linewidth=0.9)
    plt.title(f"Subject {rec} - Example 3-Second Window")
    plt.xlabel("Samples")
    plt.ylabel("Amplitude (mV)")
    plt.tight_layout()

plt.suptitle("Example ECG Windows for All Subjects", fontsize=14, y=1.02)
plt.show()
```



```
[22]: def extract_full_features(windows, fs):
    features = []
    for w in windows:
        # Basic stats
        min_val, max_val = np.min(w), np.max(w)
        amplitude = max_val - min_val
        peaks, _ = find_peaks(w, distance=fs*0.3)
        troughs, _ = find_peaks(-w, distance=fs*0.3)
        rr_interval = np.mean(np.diff(peaks))/fs if len(peaks) > 1 else np.nan
        # Statistical descriptors
        stats = [np.mean(w), np.std(w), np.median(w), skew(w), kurtosis(w)]
        features.append([min_val, max_val, amplitude, len(peaks), len(troughs),
        rr_interval] + stats)
    columns = [
    'Min', 'Max', 'Amplitude', 'Num_Peaks', 'Num_Troughs', 'RR_Interval', 'Mean', 'Std', 'Median', 'Ske
    return pd.DataFrame(features, columns=columns)

combined_features = []
for rec in subjects:
    signal = records[rec]['signal']
    windows = segment_signal(signal, fs, window_size, overlap)
    feat = extract_full_features(windows, fs)
```

```

    feat['Subject_ID'] = rec
    combined_features.append(feats)
combined_features = pd.concat(combined_features, ignore_index=True)
print(f"\nCombined feature shape: {combined_features.shape}")

```

Combined feature shape: (3610, 12)

```

[23]: fs = 365          # sampling frequency
      window_size = 3    # seconds
      overlap = 0.5      # seconds

      subjects = ['118e_6', '118e12', '118e18', '119e_6', '119e24']
      all_feature_tables = []

      for rec in subjects:
          signal = records[rec]['signal']
          windows = segment_signal(signal, fs, window_size, overlap)
          df_feat = extract_full_features(windows, fs)
          df_feat['Subject_ID'] = rec
          all_feature_tables.append(df_feat)
          print(f" Record {rec}: {df_feat.shape[0]} windows × {df_feat.shape[1]}
          ↳ features extracted")

      # Combine all subjects into one table
      combined_features = pd.concat(all_feature_tables, ignore_index=True)
      print(f"\nFinal combined feature table shape: {combined_features.shape}")

```

```

Record 118e_6: 712 windows × 12 features extracted
Record 118e12: 712 windows × 12 features extracted
Record 118e18: 712 windows × 12 features extracted
Record 119e_6: 712 windows × 12 features extracted
Record 119e24: 712 windows × 12 features extracted

```

Final combined feature table shape: (3560, 12)

```

[24]: subject_summary = combined_features.groupby("Subject_ID").agg({
      'Min': 'mean',
      'Max': 'mean',
      'Amplitude': 'mean',
      'Num_Peaks': 'mean',
      'Num_Troughs': 'mean',
      'RR_Interval': 'mean',
      'Mean': 'mean',
      'Std': 'mean',
      'Median': 'mean',
      'Skewness': 'mean',
      'Kurtosis': 'mean'

```

```
}).round(4)

# Display neatly
subject_summary = subject_summary.reset_index()
subject_summary
```

```
[24]:
```

	Subject_ID	Min	Max	Amplitude	Num_Peaks	Num_Troughs	RR_Interval	\
0	118e12	-7.1905	-3.8017	3.3888	8.0744	8.0815	0.3912	
1	118e18	-7.2253	-4.1231	3.1022	8.0211	8.1067	0.3936	
2	118e_6	-7.3338	2.7088	10.0426	8.2486	8.0716	0.3839	
3	119e24	-6.9957	-3.2877	3.7080	7.6390	7.9354	0.4164	
4	119e_6	-7.1068	2.2534	9.3602	8.0154	7.9958	0.3965	

	Mean	Std	Median	Skewness	Kurtosis
0	-5.6871	0.5076	-5.6830	-0.0918	3.0196
1	-5.8295	0.4353	-5.8054	-0.2551	3.5765
2	-3.7329	1.9329	-3.9665	0.1263	2.8280
3	-5.9126	0.5085	-5.9988	2.7574	11.9196
4	-4.0311	1.7696	-4.3016	1.9283	7.3471

```
[25]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from scipy.signal import find_peaks

fs = 360 # sampling frequency (Hz)

def analyze_and_plot_pqrst(signal, fs, subject_id, duration=3):
    """
    Detects P, Q, R, S, and T waves in an ECG signal segment,
    plots them, and returns amplitude & interval features.
    """
    segment = signal[:fs*duration]
    time_axis = np.arange(len(segment)) / fs

    # Detect R-peaks (dominant spikes)
    r_peaks, _ = find_peaks(segment, distance=fs*0.4, height=np.mean(segment)+0.
↪2*np.std(segment))

    # Initialize
    p_amp = q_amp = r_amp = s_amp = t_amp = np.nan
    pr_interval = qrs_duration = qt_interval = np.nan
    p_idx = q_idx = r_idx = s_idx = t_idx = None

    if len(r_peaks) > 0:
        r_idx = r_peaks[0]
        r_amp = segment[r_idx]
```

```

# Approximate other wave positions relative to R
p_idx = max(0, r_idx - int(0.2 * fs))
q_idx = max(0, r_idx - int(0.04 * fs))
s_idx = min(len(segment) - 1, r_idx + int(0.04 * fs))
t_idx = min(len(segment) - 1, r_idx + int(0.25 * fs))

p_amp, q_amp, s_amp, t_amp = segment[p_idx], segment[q_idx],
↪segment[s_idx], segment[t_idx]

pr_interval = (r_idx - p_idx) / fs
qrs_duration = (s_idx - q_idx) / fs
qt_interval = (t_idx - q_idx) / fs

# ---- Plot ----
plt.figure(figsize=(10, 3))
plt.plot(time_axis, segment, color='black', linewidth=1)
plt.title(f"Subject {subject_id} - P-Q-R-S-T Wave Detection")
plt.xlabel("Time (seconds)")
plt.ylabel("Amplitude (mV)")
plt.grid(alpha=0.3)

plt.scatter(time_axis[[p_idx, q_idx, r_idx, s_idx, t_idx]],
            segment[[p_idx, q_idx, r_idx, s_idx, t_idx]],
            color=['blue', 'orange', 'red', 'green', 'purple'],
            s=50, zorder=3)

labels = ['P', 'Q', 'R', 'S', 'T']
for idx, label, color in zip([p_idx, q_idx, r_idx, s_idx, t_idx],
                             labels,
                             ['blue', 'orange', 'red', 'green',
↪'purple']):
    plt.text(time_axis[idx], segment[idx] + 0.05, label,
             color=color, fontsize=9, ha='center', fontweight='bold')
plt.show()

# Return numerical features
return {
    'Subject_ID': subject_id,
    'P_amp': p_amp, 'Q_amp': q_amp, 'R_amp': r_amp, 'S_amp': s_amp, 'T_amp':
↪t_amp,
    'PR_interval': pr_interval, 'QRS_duration': qrs_duration, 'QT_interval':
↪qt_interval
}

# -----

```


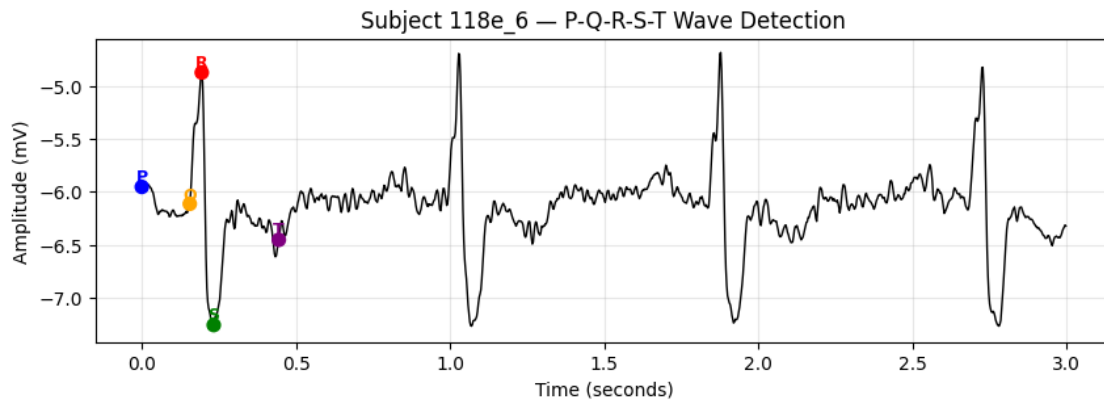
```


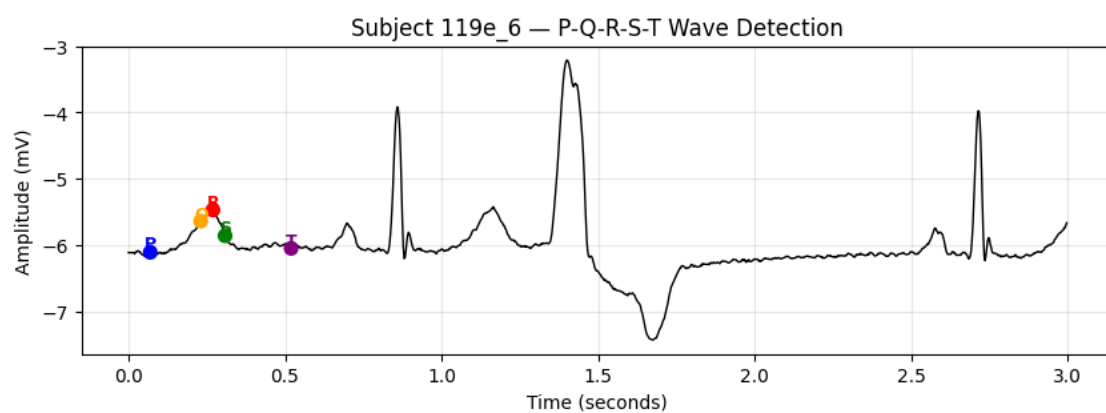

# Run for all five selected subjects and collect data
# -----
subjects = ['118e_6', '118e12', '118e18', '119e_6', '119e24']
results = []

for rec in subjects:
    signal = records[rec]['signal']
    features = analyze_and_plot_pqrst(signal, fs, rec, duration=3)
    results.append(features)

# Convert to DataFrame for a clean summary
df_pqrst_summary = pd.DataFrame(results)
print("\nExtracted PQRST Amplitude and Interval Features for 5 Subjects:\n")
print(df_pqrst_summary.round(4))

```





Extracted PQRST Amplitude and Interval Features for 5 Subjects:

	Subject_ID	P_amp	Q_amp	R_amp	S_amp	T_amp	PR_interval	QRS_duration	\
0	118e_6	-5.955	-6.115	-4.875	-7.26	-6.455	0.1944	0.0778	
1	118e12	-5.955	-6.115	-4.875	-7.26	-6.455	0.1944	0.0778	
2	118e18	-5.955	-6.115	-4.875	-7.26	-6.455	0.1944	0.0778	
3	119e_6	-6.110	-5.640	-5.470	-5.86	-6.050	0.2000	0.0778	
4	119e24	-6.110	-5.640	-5.470	-5.86	-6.050	0.2000	0.0778	

	QT_interval
0	0.2889
1	0.2889
2	0.2889
3	0.2889
4	0.2889

1 Step 3

```
[27]: try:
      df
    except NameError:
      df = combined_features.copy()

merged_list = []

for subj in df["Subject_ID"].unique():
    feat_sub = df[df["Subject_ID"] == subj].copy().reset_index(drop=True)
    label_sub = target_df[target_df["Subject_ID"] == subj].copy().
    ↪reset_index(drop=True)

    # Align window counts
    min_len = min(len(feat_sub), len(label_sub))
    feat_sub = feat_sub.iloc[:min_len]
    label_sub = label_sub.iloc[:min_len]

    feat_sub["Target"] = label_sub["Target"]
    merged_list.append(feat_sub)

# Combine all subjects
df_final = pd.concat(merged_list, ignore_index=True)

print(" Final dataset created successfully!")
print("Shape:", df_final.shape)
print(df_final["Target"].value_counts())
```

```
Final dataset created successfully!
Shape: (3560, 13)
Target
0      2201
```

```
1    1359
Name: count, dtype: int64
```

```
[28]: X = df_final.drop(columns=["Subject_ID", "Target"])
      y = df_final["Target"]

      # Balance dataset using SMOTE
      sm = SMOTE(random_state=42)
      X_res, y_res = sm.fit_resample(X, y)
      print("After SMOTE balancing:")
      print(pd.Series(y_res).value_counts())

      # Split train/test
      X_train, X_test, y_train, y_test = train_test_split(
          X_res, y_res, test_size=0.2, random_state=42, stratify=y_res
      )

      # Scale features
      scaler = StandardScaler()
      X_train_scaled = scaler.fit_transform(X_train)
      X_test_scaled = scaler.transform(X_test)
      print(" Data prepared and scaled.")
```

```
After SMOTE balancing:
Target
1    2201
0    2201
Name: count, dtype: int64
 Data prepared and scaled.
```

```
[29]: # Define models
      slnn = MLPClassifier(hidden_layer_sizes=(10,), activation='relu',
                           solver='adam', max_iter=500, random_state=42)

      mlp = MLPClassifier(hidden_layer_sizes=(32,16), activation='relu',
                           solver='adam', learning_rate_init=0.001,
                           early_stopping=True, max_iter=1000, random_state=42)

      # Train models
      slnn.fit(X_train_scaled, y_train)
      mlp.fit(X_train_scaled, y_train)

      print("Models trained successfully!")
```

```
Models trained successfully!
```

```
[36]: subject_results = []
```

```

plt.figure(figsize=(7,5))

for subj in subjects:
    # Filter data for this subject and remove missing targets
    sub_df = df_final[df_final['Subject_ID'] == subj].dropna(subset=['Target'])

    if sub_df.empty:
        print(f"Skipping {subj} - no valid Target data.")
        continue

    # Extract features and scale
    X_raw = sub_df.drop(columns=['Subject_ID', 'Target'])
    X_sub = scaler.transform(X_raw)
    y_sub = sub_df['Target'].astype(int)

    # Clean scaled data: handle NaN, inf, and extreme values
    X_sub = np.nan_to_num(X_sub, nan=0.0, posinf=0.0, neginf=0.0)
    imputer = SimpleImputer(strategy='mean')
    X_sub = imputer.fit_transform(X_sub)
    X_sub = np.clip(X_sub, -1e6, 1e6)

    try:
        # Predictions
        y_pred = mlp.predict(X_sub)
        y_prob = mlp.predict_proba(X_sub)[:, 1]

        # Compute metrics
        acc = accuracy_score(y_sub, y_pred)
        f1 = f1_score(y_sub, y_pred)
        auc_val = roc_auc_score(y_sub, y_prob)

        subject_results.append({
            'Subject_ID': subj,
            'Accuracy': round(acc, 3),
            'F1_Score': round(f1, 3),
            'AUC': round(auc_val, 3)
        })

        # Plot ROC Curve for this subject
        fpr, tpr, _ = roc_curve(y_sub, y_prob)
        plt.plot(fpr, tpr, lw=1.5, label=f"{subj} (AUC={auc_val:.2f})")

    except ValueError as e:
        print(f"Skipping {subj} due to error: {e}")

# Combined ROC visualization
plt.plot([0,1], [0,1], '--', color='gray')

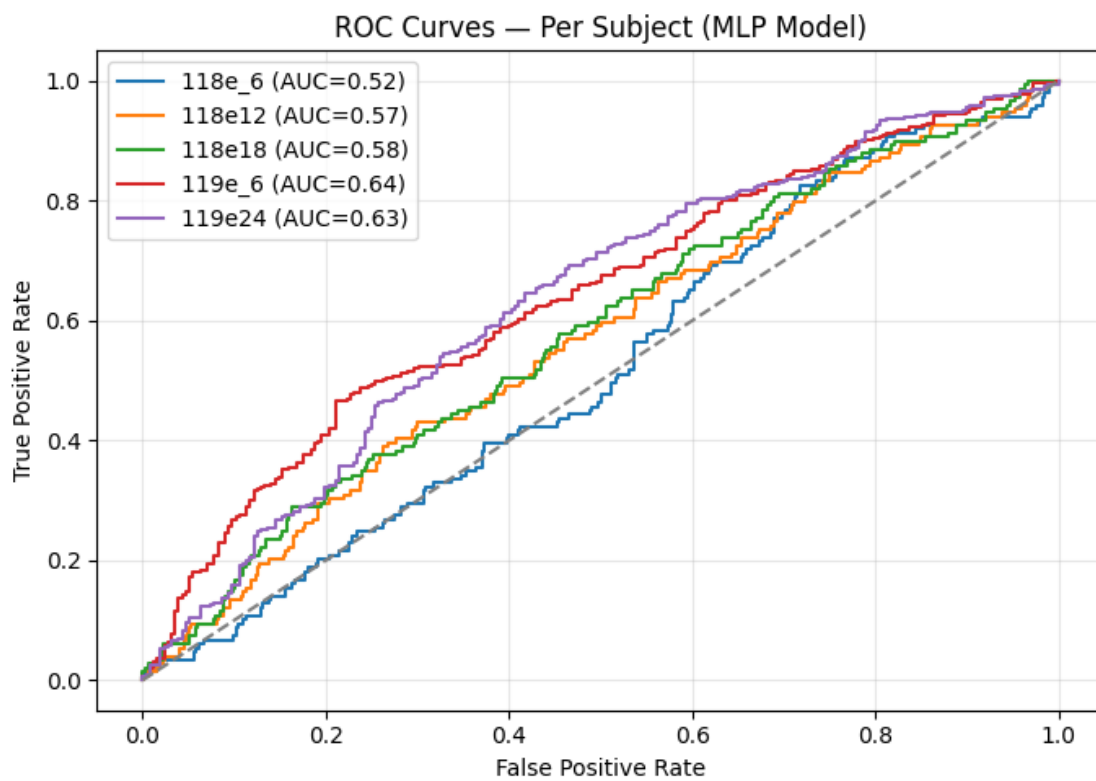
```

```

plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ROC Curves - Per Subject (MLP Model)")
plt.legend()
plt.grid(alpha=0.3)
plt.tight_layout()
plt.show()

# Display results summary
results_df = pd.DataFrame(subject_results)
print("\n Per-Subject Evaluation Results:\n")
display(results_df)

```



Per-Subject Evaluation Results:

	Subject_ID	Accuracy	F1_Score	AUC
0	118e_6	0.687	0.189	0.520
1	118e12	0.779	0.071	0.567
2	118e18	0.784	0.061	0.580
3	119e_6	0.638	0.734	0.642
4	119e24	0.640	0.779	0.633

```
[30]: results = []

for name, model in [("Single-Layer NN", slnn), ("Multi-Layer Perceptron", mlp)]:
    y_pred = model.predict(X_test_scaled)
    y_prob = model.predict_proba(X_test_scaled)[: , 1]

    metrics = {
        "Model": name,
        "Accuracy": accuracy_score(y_test, y_pred),
        "Precision": precision_score(y_test, y_pred),
        "Recall": recall_score(y_test, y_pred),
        "F1": f1_score(y_test, y_pred),
        "AUC": roc_auc_score(y_test, y_prob)
    }
    results.append(metrics)

    print(f"\n{name} Classification Report:\n", classification_report(y_test,
↪y_pred))

# Display summary table
results_df = pd.DataFrame(results).round(3)
print("\n=== MODEL COMPARISON ===")
print(results_df)
```

Single-Layer NN Classification Report:

	precision	recall	f1-score	support
0	0.68	0.73	0.70	441
1	0.71	0.65	0.68	440
accuracy			0.69	881
macro avg	0.69	0.69	0.69	881
weighted avg	0.69	0.69	0.69	881

Multi-Layer Perceptron Classification Report:

	precision	recall	f1-score	support
0	0.65	0.73	0.69	441
1	0.69	0.61	0.65	440
accuracy			0.67	881
macro avg	0.67	0.67	0.67	881
weighted avg	0.67	0.67	0.67	881

=== MODEL COMPARISON ===

	Model	Accuracy	Precision	Recall	F1	AUC
0	Single-Layer NN	0.691	0.708	0.650	0.678	0.745
1	Multi-Layer Perceptron	0.670	0.692	0.611	0.649	0.725

```
[33]: subject_results_slmn = []
      subject_results_mlp = []

      plt.figure(figsize=(8,6))
      plt.title("ROC Curve per Subject - Single-Layer NN")
      for subj in df_final["Subject_ID"].unique():
          sub_df = df_final[df_final["Subject_ID"] == subj].copy().
          ↪dropna(subset=["Target"])

          # Prepare data
          X_sub = sub_df.drop(columns=["Subject_ID", "Target"])
          y_sub = sub_df["Target"].astype(int)
          X_sub_scaled = scaler.transform(X_sub)


          # Predictions
          y_prob = slmn.predict_proba(X_sub_scaled)[: , 1]
          y_pred = slmn.predict(X_sub_scaled)

          # Metrics
          acc = accuracy_score(y_sub, y_pred)
          f1 = f1_score(y_sub, y_pred)
          auc_val = roc_auc_score(y_sub, y_prob)
          subject_results_slmn.append({"Subject_ID": subj, "Accuracy": acc, "F1": f1,
          ↪"AUC": auc_val})

          # ROC curve
          fpr, tpr, _ = roc_curve(y_sub, y_prob)
          plt.plot(fpr, tpr, lw=1.5, label=f"{subj} (AUC={auc_val:.2f})")

      plt.plot([0,1],[0,1], '--', color='gray')
      plt.xlabel("False Positive Rate")
      plt.ylabel("True Positive Rate")
      plt.legend()
      plt.grid(alpha=0.3)
      plt.show()

      # Convert results to DataFrame
      subject_results_slmn = pd.DataFrame(subject_results_slmn)
      print("\n=== Subject-wise Performance: Single-Layer NN ===")
      print(subject_results_slmn.round(3))
```



=== Subject-wise Performance: Single-Layer NN ===

	Subject_ID	Accuracy	F1	AUC
0	118e_6	0.687	0.228	0.569
1	118e12	0.770	0.118	0.606
2	118e18	0.774	0.101	0.619
3	119e_6	0.660	0.758	0.667
4	119e24	0.643	0.778	0.636

```
[34]: plt.figure(figsize=(8,6))
plt.title("ROC Curve per Subject - Multi-Layer Perceptron")
subject_results_mlp = []

for subj in df_final["Subject_ID"].unique():
    sub_df = df_final[df_final["Subject_ID"] == subj].copy().
    dropna(subset=["Target"])

    X_sub = sub_df.drop(columns=["Subject_ID", "Target"])
    y_sub = sub_df["Target"].astype(int)
```

```

X_sub_scaled = scaler.transform(X_sub)

y_prob = mlp.predict_proba(X_sub_scaled)[: , 1]
y_pred = mlp.predict(X_sub_scaled)


acc = accuracy_score(y_sub, y_pred)
f1 = f1_score(y_sub, y_pred)
auc_val = roc_auc_score(y_sub, y_prob)
subject_results_mlp.append({"Subject_ID": subj, "Accuracy": acc, "F1": f1,
↪ "AUC": auc_val})

fpr, tpr, _ = roc_curve(y_sub, y_prob)
plt.plot(fpr, tpr, lw=1.5, label=f"{subj} (AUC={auc_val:.2f})")

plt.plot([0,1],[0,1], '--', color='gray')
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.legend()
plt.grid(alpha=0.3)
plt.show()

# Convert to DataFrame
subject_results_mlp = pd.DataFrame(subject_results_mlp)
print("\n=== Subject-wise Performance: Multi-Layer Perceptron ===")
print(subject_results_mlp.round(3))

```

=== Subject-wise Performance: Multi-Layer Perceptron ===

	Subject_ID	Accuracy	F1	AUC
0	118e_6	0.687	0.189	0.520
1	118e12	0.779	0.071	0.567
2	118e18	0.784	0.061	0.580
3	119e_6	0.638	0.734	0.642
4	119e24	0.640	0.779	0.633

2 Optional

```
[35]: from sklearn.neural_network import MLPClassifier
      from sklearn.metrics import accuracy_score, f1_score, roc_auc_score

      # Features and target
      X = df.drop(columns=["Subject_ID", "Target"])
      y = df["Target"]

      # Train-test split
```

```

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y
)

# Scale
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Single-Layer NN (SLNN)
slnn = MLPClassifier(hidden_layer_sizes=(10,), activation='relu',
    ↪solver='adam', max_iter=500, random_state=42)
slnn.fit(X_train_scaled, y_train)
y_pred_slmn = slnn.predict(X_test_scaled)
y_prob_slmn = slnn.predict_proba(X_test_scaled)[: ,1]

# Multi-Layer Perceptron (MLP)
mlp = MLPClassifier(hidden_layer_sizes=(32,16), activation='relu',
    ↪solver='adam', max_iter=700, random_state=42)
mlp.fit(X_train_scaled, y_train)
y_pred_mlp = mlp.predict(X_test_scaled)
y_prob_mlp = mlp.predict_proba(X_test_scaled)[: ,1]

# Compare performance
comparison = pd.DataFrame({
    "Model": ["Single-Layer NN", "Multi-Layer Perceptron"],
    "Accuracy": [accuracy_score(y_test, y_pred_slmn), accuracy_score(y_test,
    ↪y_pred_mlp)],
    "F1": [f1_score(y_test, y_pred_slmn), f1_score(y_test, y_pred_mlp)],
    "AUC": [roc_auc_score(y_test, y_prob_slmn), roc_auc_score(y_test,
    ↪y_prob_mlp)]
}).round(3)

comparison

```

```

[35]:


```

	Model	Accuracy	F1	AUC
0	Single-Layer NN	0.501	0.482	0.475
1	Multi-Layer Perceptron	0.475	0.416	0.473

```

[36]: # Feature Importance and Correlation Visualization
plt.figure(figsize=(10,8))
corr = X.corr()
sns.heatmap(corr, cmap='coolwarm', annot=False)
plt.title("Feature Correlation Heatmap")
plt.show()

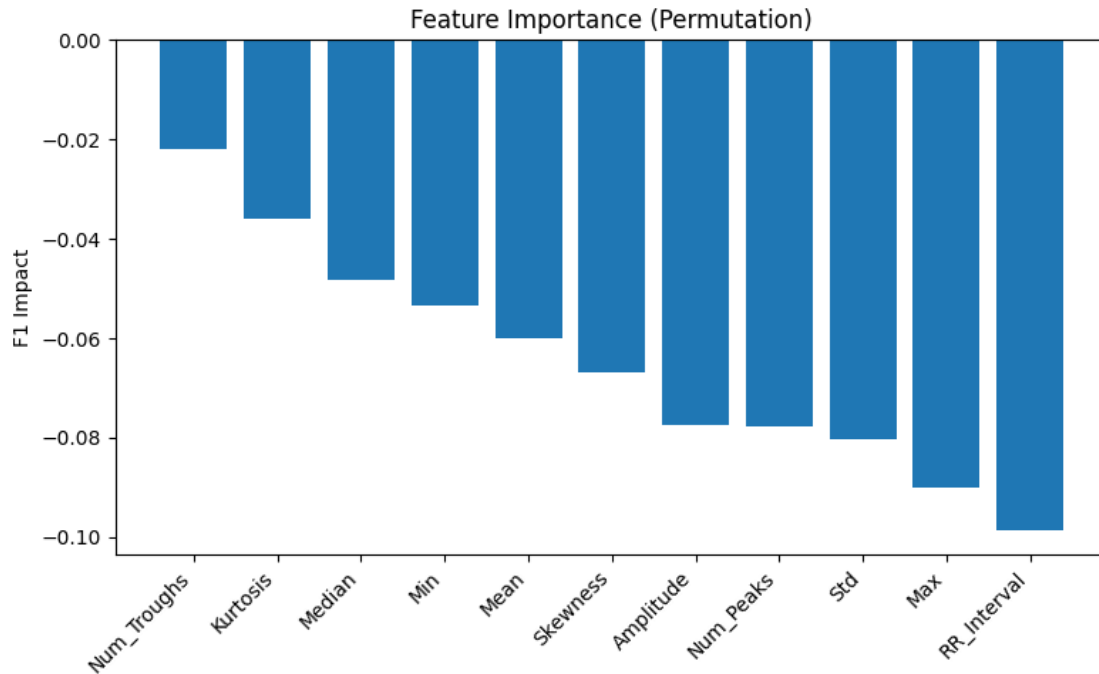
```



```
[37]: # Permutation Importance
from sklearn.inspection import permutation_importance

perm_importance = permutation_importance(mlp, X_test_scaled, y_test,
    ↪scoring='f1', n_repeats=10, random_state=42)

sorted_idx = perm_importance.importances_mean.argsort()[::-1]
plt.figure(figsize=(8,5))
plt.bar(range(len(sorted_idx)), perm_importance.importances_mean[sorted_idx])
plt.xticks(range(len(sorted_idx)), X.columns[sorted_idx], rotation=45,
    ↪ha='right')
plt.title("Feature Importance (Permutation)")
plt.ylabel("F1 Impact")
plt.tight_layout()
plt.show()
```



```
[38]: # Real time window prediction
fs = 360
signal = records['118e_6']['signal']

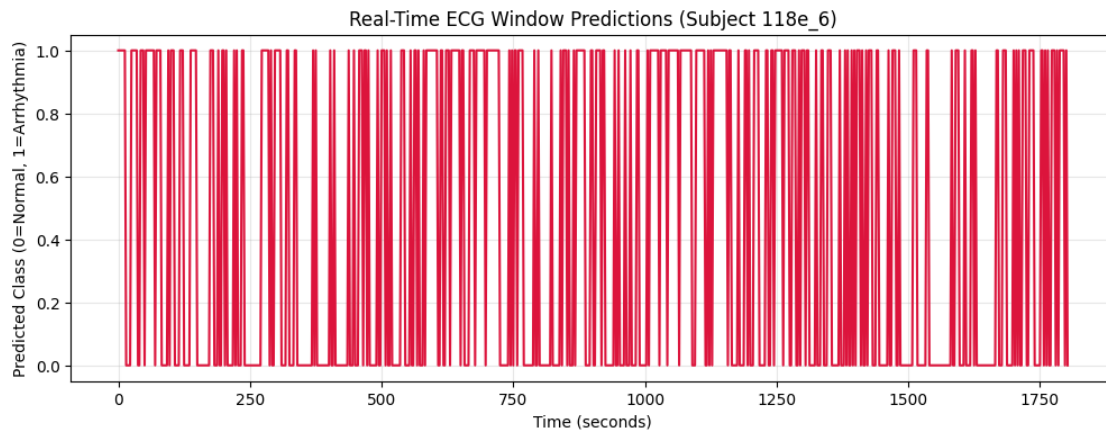
# Sliding window params
window_size = 3 # seconds
overlap = 0.5 # seconds
win_len = int(fs * window_size)
step = int(fs * (window_size - overlap))

window_preds = []
time_points = []

for start in range(0, len(signal) - win_len, step):
    window = signal[start:start+win_len]
    # extract features same way as before
    feat = extract_full_features([window], fs)
    X_window = scaler.transform(feat[X.columns]) # same scaling as training
    pred = mlp.predict(X_window)[0]
    window_preds.append(pred)
    time_points.append(start / fs) # seconds

# Convert to DataFrame for plotting
df_real_time = pd.DataFrame({'Time_s': time_points, 'Prediction': window_preds})
```

```
plt.figure(figsize=(12,4))
plt.plot(df_real_time['Time_s'], df_real_time['Prediction'], color='crimson', lw=1.5)
plt.title("Real-Time ECG Window Predictions (Subject 118e_6)")
plt.xlabel("Time (seconds)")
plt.ylabel("Predicted Class (0=Normal, 1=Arrhythmia)")
plt.grid(alpha=0.3)
plt.show()
```



[]: